Use the `head` command on your three files again. This time, describe at least one potential problem with the data you see. Consider issues with missing values and bad data.

```
In [17]: bus.head(50)
         #ins2vio.head()
         #ins.head()
         #vio.head()
```

Out[17]:

	business id column	name
0	1000	HEUNG YUEN RESTAURANT
1	100010	ILLY CAFFE SF_PIER 39
2	100017	AMICI'S EAST COAST PIZZERIA
3	100026	LOCAL CATERING
4	100030	OUI OUI! MACARON
5	100036	Hula Truck (#2)
6	100039	GENKI CREPES & MINI MART
7	100041	UNCLE LEE CAFE
8	100055	Twirl and Dip
9	100058	SF PITA HUB
10	100059	DUMPLING ALLEY
11	100069	Mission Blue
12	100072	SUBWAY SANDWICHES #7307
13	100079	POSITIVE FOODS
14	100081	THE MATTERHORN RESTAURANT AND BAKERY
15	100082	SLN CTRNG
16	100083	THE EPICUREAN TRADER
17	100084	FRJTZ KITCHEN
18	100096	THE LITTLE CHIHUAHUA MEXICAN RESTAURANT
19	100097	GANGNAM BBQ
20	100098	ZHONG SHAN RESTAURANT
21	100099	KEN KEE CAFE
22	100126	Lamas Peruvian Food Truck
23	100135	Hotel Whitcomb - Employee Kitchen
24	100137	100137 Cloud Club
25	100142	Multi Service Center South
26	100145	Conchinita
27	1002	BIG MOUTH BURGERS
28	100202	FACEBOOK INC.
29	100203	FACEBOOK, INC.
30	100204	CUIA ACAI & POSITIVE FOOD
31	100205	HUNAN EMPIRE
32	100210	KING OF THAI NOODLE HOUSE
33	100211	THE EAGLE CAFE
34	100212	BELCAMPO MEAT CO
35	100214	REAL KABOB
36	100215	CHICKEN N WAFFLES PLACE
37	100216	BUNN MIKE
38	100219	NAYA CAFE
39	100238	PEACHES PATTIES LLC
40	100239	HUNTINGTON HOTEL SAN FRANCISCO
41	100240	U :DESSERT STORY

42	100241	Taqueria San Marcos				
43	100252	BITE ME SANDWICHES				
44	100253	BISTRO LOVESSY, LLC				
45	100255	H&M FOOD MART				
46	100274	THE OLYMPIC CAFE				
47	100275	LITTLE CREATURES BREWING COMPANY				
48	100277	COAST TO COAST ACAI AND GRANOLA				
49	100278	KINARA KITCHEN, INC.				

	address	city	state	postal_code	\
0	3279 22nd St	San Francisco	CA	94110	
1	PIER 39 K-106-B	San Francisco	CA	94133	
2	475 06th St	San Francisco	CA	94103	
3	1566 CARROLL AVE	San Francisco	CA	94124	
4	2200 JERROLD AVE STE C	San Francisco	CA	94124	
5	2 Marina Blvd	San Francisco	CA	94123	
6	330 CLEMENT ST	San Francisco	CA	94118	
7	3608 BALBOA ST	San Francisco	CA	94121	
8	335 Martin Luther King Jr. Dr	San Francisco	CA	94118	
9	475 06TH ST	San Francisco	CA	94103	
10	2512 CLEMENT ST	San Francisco	CA	94121	
11	144 Leland Ave	San Francisco	CA	94134	
12	2375 MARKET ST	San Francisco	CA	94114	
13	475 06TH ST	San Francisco	CA	94103	
14	2323 VAN NESS AVE	San Francisco	CA	94109	
15	103 HORNE Ave	San Francisco	CA	94124	
16	465 HAYES ST	San Francisco	CA	94102	
17	475 06TH ST UNIT 15	San Francisco	CA	94103	
18	475 06TH ST K16	San Francisco	CA	94103	
19	3251 20TH AVE 250B	San Francisco	CA	94132	
20	2237 TARAVAL ST	San Francisco	CA	94116	
21	2109 CLEMENT ST	San Francisco	CA	94121	
22	Private Location	San Francisco	CA	-9999	
23	1231 Market St	San Francisco	CA	94103	
24	24 Willie Mays Plaza Suites Level	San Francisco	CA	94107	
25	525 05th St	San Francisco	CA	94107	
26	2 Marina Blvd Fort Mason	San Francisco	CA	94123	
27	3392 24th St	San Francisco	CA	94110	
28	181 FREMONT ST FL 5TH	San Francisco	CA	94105	
29	181 FREMONT ST FL 6TH	San Francisco	CA	94105	
30	1 MARKET ST STE 8	San Francisco	CA	94105	
31	2001 UNION ST #107	San Francisco	CA	94123	
32	184 O'FARRELL ST	San Francisco	CA	94102	
33	39 PIER A201	San Francisco	CA	94133	
34	475 06TH ST	San Francisco	CA	94103	
35	475 06TH St 23	San Francisco	CA	94103	
36	1968 LOMBARD ST	San Francisco	CA	94123	
37	300 DE HARO ST	San Francisco	CA	94103	
38	5338 geary BLVD	San Francisco	CA	94121	
39	2948 FOLSOM ST	San Francisco	CA	94110	
40	1075 CALIFORNIA ST	San Francisco	CA	94108	
41	2120 GREENWICH ST	San Francisco	CA	94123	
42	2380 San Bruno Ave	San Francisco	CA	94134	
43	701 COLE ST	San Francisco	CA	94117	

44	832 CLEMENT ST	San Francisco	CA	94118
45	2400 SAN BRUNO AVE	San Francisco	CA	94134
46	555 GEARY ST	San Francisco	CA	94102
47	1000 A 03RD St	San Francisco	CA	94158
48	160 14TH STREET	San Francisco	CA	94103
49	607 GEARY ST	San Francisco	CA	94102

	latitude	longitude	phone_number
0	37.755282	-122.420493	-9999
1	-9999.000000	-9999.000000	14154827284
2	-9999.000000	-9999.000000	14155279839
3	-9999.000000	-9999.000000	14155860315
4	-9999.000000	-9999.000000	14159702675
5	-9999.000000	-9999.000000	-9999
6	-9999.000000	-9999.000000	14155376414
7	-9999.000000	-9999.000000	-9999
8	-9999.000000	-9999.000000	14155300260
9	-9999.000000	-9999.000000	14155642006
10	-9999.000000	-9999.000000	-9999
11	-9999.000000	-9999.000000	-9999
12	-9999.000000	-9999.000000	14155981866
13	-9999.000000	-9999.000000	14155397209
14	-9999.000000	-9999.000000	14155474029
15	-9999.000000	-9999.000000	14155965620
16	-9999.000000	-9999.000000	14155606092
17	-9999.000000	-9999.000000	14155868272
18	-9999.000000	-9999.000000	-9999
19	-9999.000000	-9999.000000	14150494183
20	-9999.000000	-9999.000000	14155806898
21	-9999.000000	-9999.000000	14155699118
22	-9999.000000	-9999.000000	-9999
23	-9999.000000	-9999.000000	-9999
24	-9999.000000	-9999.000000	-9999
25	-9999.000000	-9999.000000	-9999
26	-9999.000000	-9999.000000	-9999
27	37.752158	-122.420362	-9999
28	-9999.000000	-9999.000000	14150799045
29	-9999.000000	-9999.000000	14150799045
30	-9999.000000	-9999.000000	14158609815
31	-9999.000000	-9999.000000	14155774735
32	-9999.000000	-9999.000000	14155821999
33	-9999.000000	-9999.000000	14155985872
34	-9999.000000	-9999.000000	14157800656
35	-9999.000000	-9999.000000	14158705851
36	-9999.000000	-9999.000000	14156425140
37	-9999.000000	-9999.000000	14155299775
38	-9999.000000	-9999.000000	14155995527
39	-9999.000000	-9999.000000	-9999
40	-9999.000000	-9999.000000	14155342803
41	-9999.000000	-9999.000000	14155333435
42	-9999.000000	-9999.000000	-9999
43	-9999.000000	-9999.000000	14155665282
44	-9999.000000	-9999.000000	14155827593
45	-9999.000000	-9999.000000	14159277470

```
46 -9999.000000 -9999.000000   14155718182
47 -9999.000000 -9999.000000   14153334433
48 -9999.000000 -9999.000000   14159230622
49 -9999.000000 -9999.000000         -9999
```

- In the bus dataframe we see that the address column has different fromatting. For example, 475 06th St vs 1566 CARROLL AVE
- The phone number column also has missing values denoted as -9999
- the latitude and longitude columns are missing

In the cell below, write the name of the restaurant with the lowest inspection scores ever. You can also head to yelp.com and look up the reviews page for this restaurant. Feel free to add anything interesting you want to share.

'Lollipot" this is true because it shows the WORST score. I could not find this on yelp, maybe it was closed because the restaurant had the worst score.

0.1 Question 6a

Let's look at the distribution of inspection scores. As we saw before when we called head on this data frame, inspection scores appear to be integer values. The discreteness of this variable means that we can use a barplot to visualize the distribution of the inspection score. Make a bar plot of the counts of the number of inspections receiving each score.

It should look like the image below. It does not need to look exactly the same (e.g., no grid), but make sure that all labels and axes are correct.

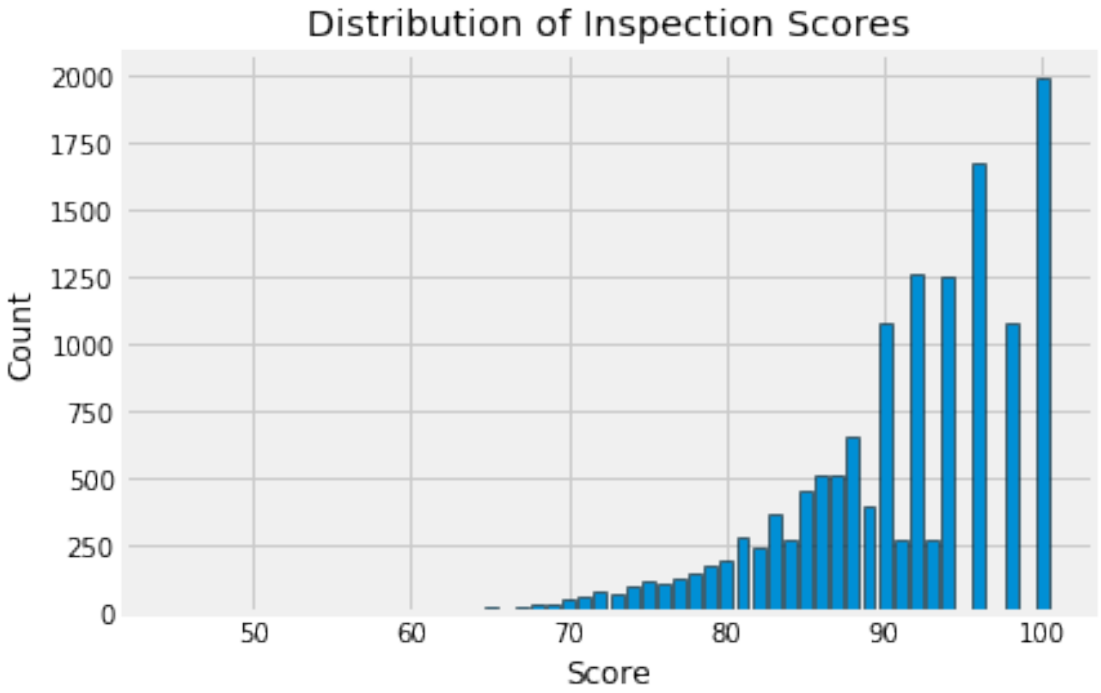

You might find this matplotlib.pyplot tutorial useful. Key syntax that you'll need:

```
plt.bar
plt.xlabel
plt.ylabel
plt.title
```

Note: If you want to use another plotting library for your plots (e.g. plotly, sns) you are welcome to use that library instead so long as it works on DataHub. If you use seaborn sns.countplot(), you may need to manually set what to display on xticks.

In [78]:
```
x = ins['score'].value_counts().keys()
y = ins['score'].value_counts()
plt.bar(x,y)
plt.xlabel('Score')
plt.ylabel('Count')
plt.title('Distribution of Inspection Scores')
```

Out[78]: Text(0.5, 1.0, 'Distribution of Inspection Scores')

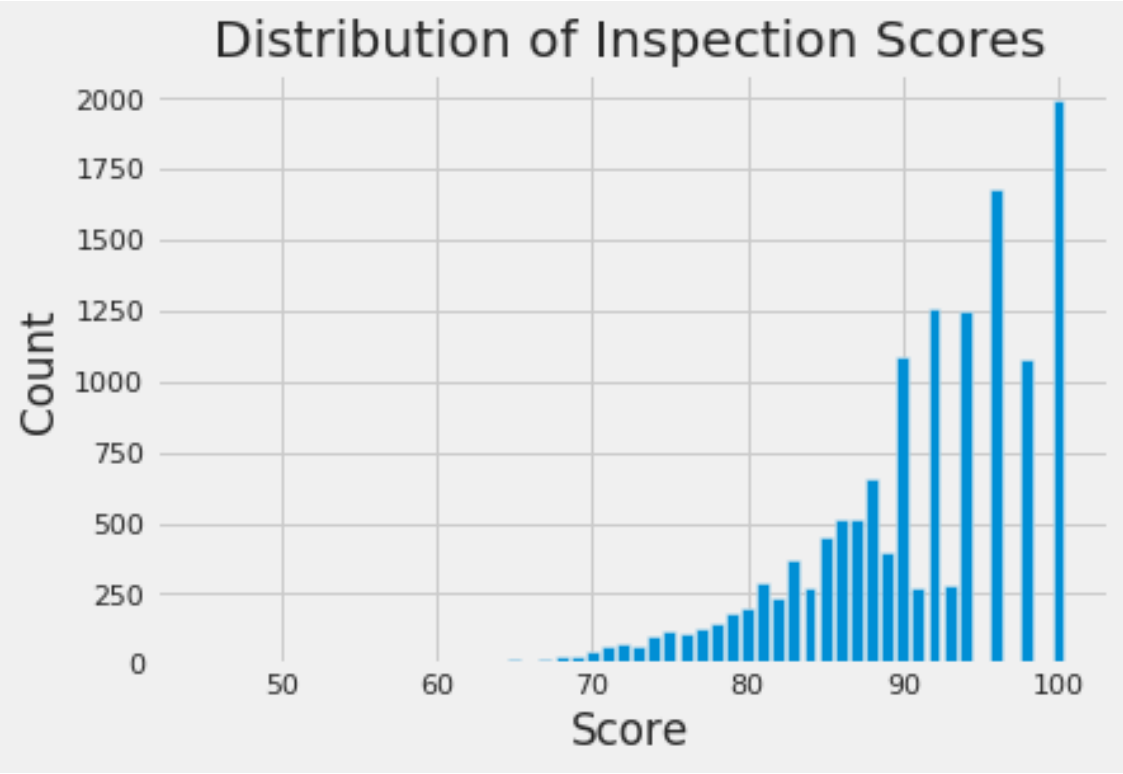

0.1.1 Question 6b

Describe the qualities of the distribution of the inspections scores based on your bar plot. Consider the mode(s), symmetry, tails, gaps, and anomalous values. Are there any unusual features of this distribution? What do your observations imply about the scores?

- The mode is the most frequent score and in the distribution above, the most frequent score is 100. We also have some other high scores that include 97,95,93,91.
- There is no evidence of symmetry in this graph. Instead this is a negatively skewed graph, the peak is on the right side with a relatively long left negative tail.
- We see gaps in three places: x=94, 96, 98.
- It seems that there is no scores that equal to the gaps of 94, 96, 98 which makes me wonder how the subcategories are calculated. Maybe there is no way to get those scores mathematically. I am pleasantly surpised that the highest scores are 100/100. This means that most restaurants are following the insepction codes and also the other majority of the scores seem to be in the upper 90s in terms of score out of 100.

Now, create your scatter plot in the cell below. It does not need to look exactly the same (e.g., no grid) as the sample below, but make sure that all labels, axes and data itself are correct.

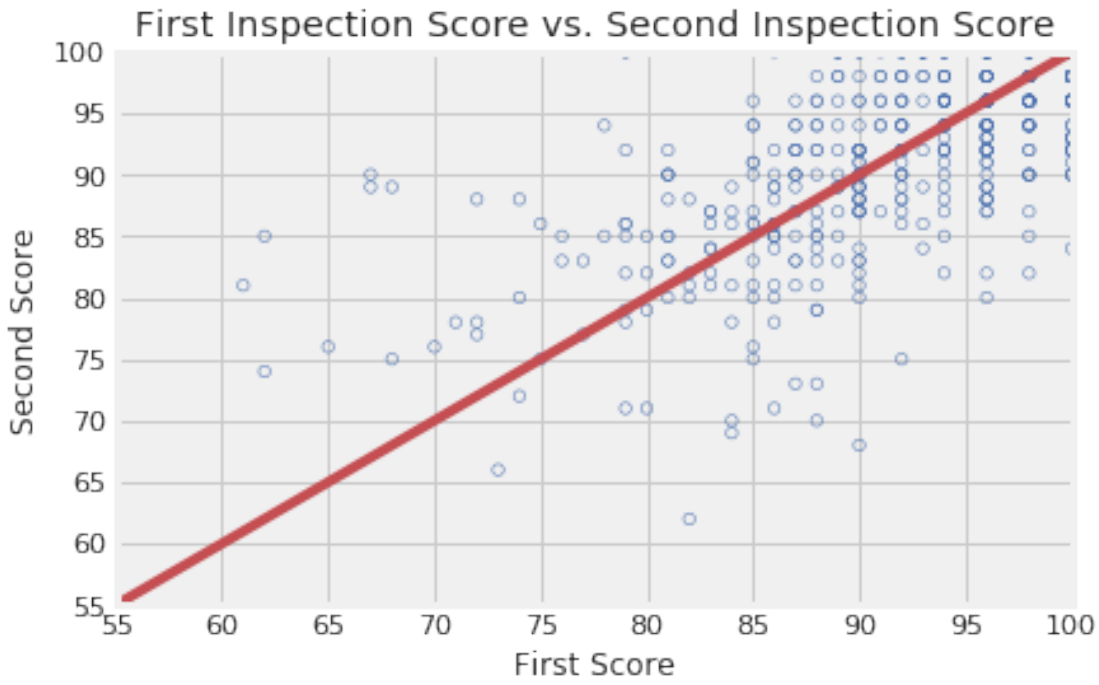

Key pieces of syntax you'll need:

`plt.scatter` plots a set of points. Use `facecolors='none'` and `edgecolors=b` to make circle markers with blue borders.

`plt.plot` for the reference line.

`plt.xlabel`, `plt.ylabel`, `plt.axis`, and `plt.title`.

Hint: You may find it convenient to use the `zip()` function to unzip scores in the list.

```
In [87]: x = scores_pairs_by_business['score_pair'].agg(lambda x: x[0]).to_list()
         y = scores_pairs_by_business['score_pair'].agg(lambda x: x[1]).to_list()
         plt.scatter(x, y, facecolors = 'none', edgecolors='b')
         plt.axis([55, 100, 55, 100])
         plt.plot([55,100],[55,100], c='r')
         plt.xlabel('First Score')
         plt.ylabel('Second Score')
         plt.title('First Inspection Score vs. Second Inspection Score');
```

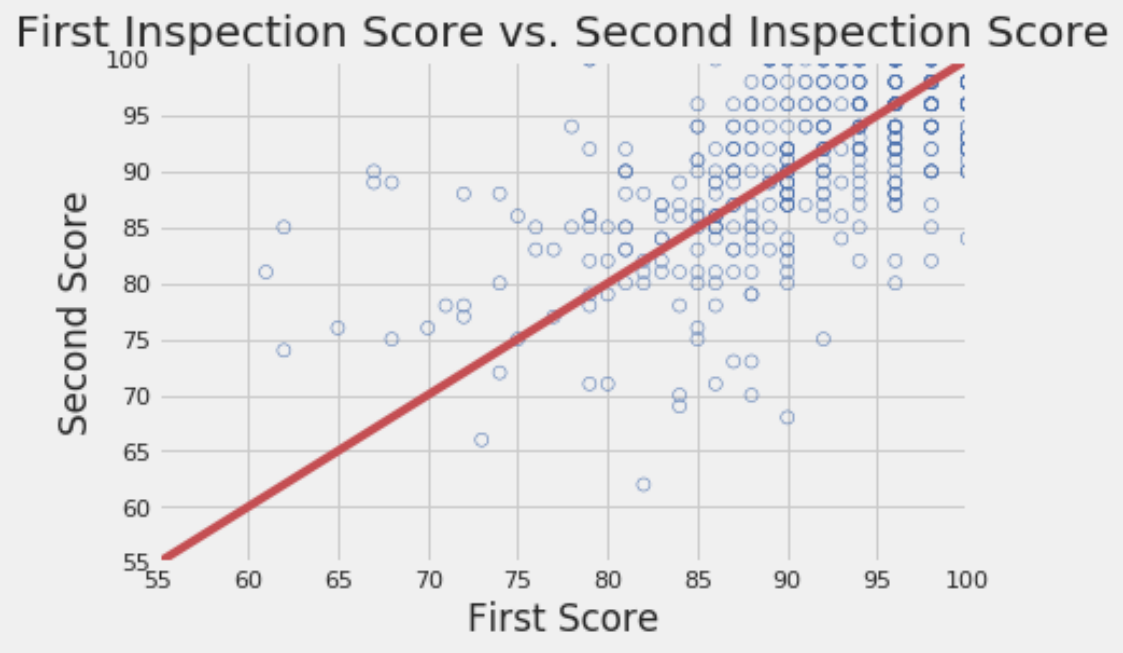
First Inspection Score vs. Second Inspection Score
Second Score
First Score
100
95
90
85
80
75
70
65
60
55
55
60
65
70
75
80
85
90
95
100

0.1.2 Question 7d

Another way to compare the scores from the two inspections is to examine the difference in scores. Subtract the first score from the second in `scores_pairs_by_business`. Make a histogram of these differences in the scores. We might expect these differences to be positive, indicating an improvement from the first to the second inspection.

The histogram should look like this:

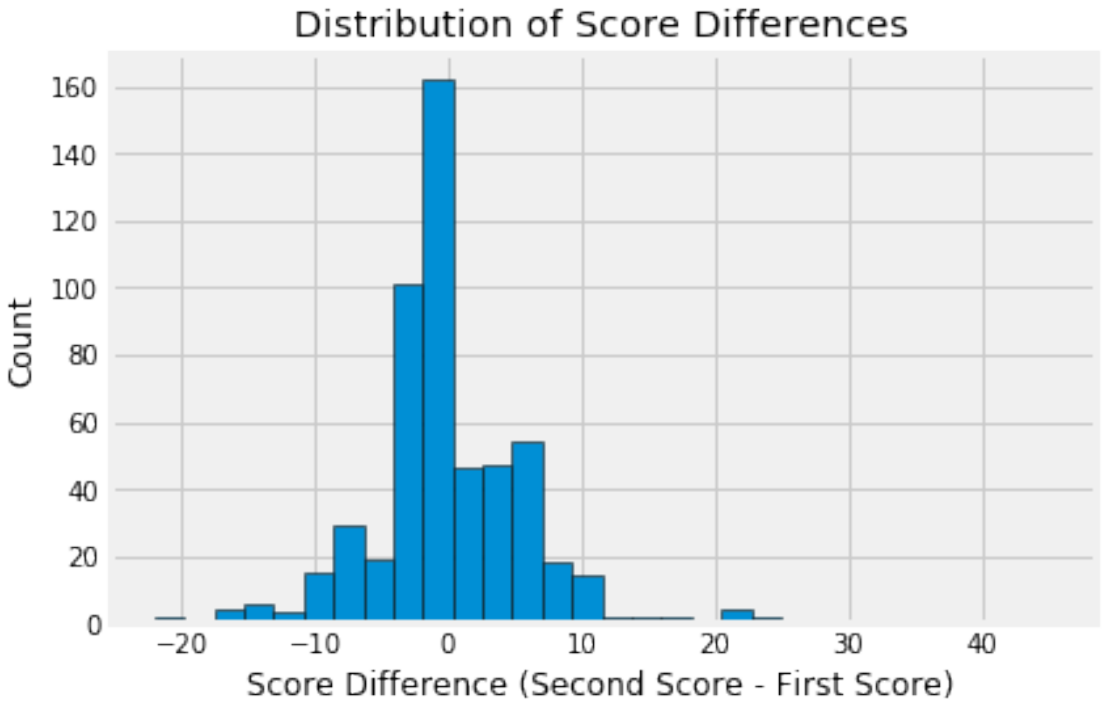

Hint: Use `second_score` and `first_score` created in the scatter plot code above.

Hint: Convert the scores into numpy arrays to make them easier to deal with.

Hint: Use `plt.hist()` Try changing the number of bins when you call `plt.hist()`.

```
In [88]: diff = [score_pair[1] - score_pair[0] for t in scores_pairs_by_business.values for score_pair
         plt.hist(diff, bins=30);
         plt.xlabel('Score Difference (Second Score - First Score)')
         plt.ylabel('Count')
         plt.title('Distribution of Score Differences');
```

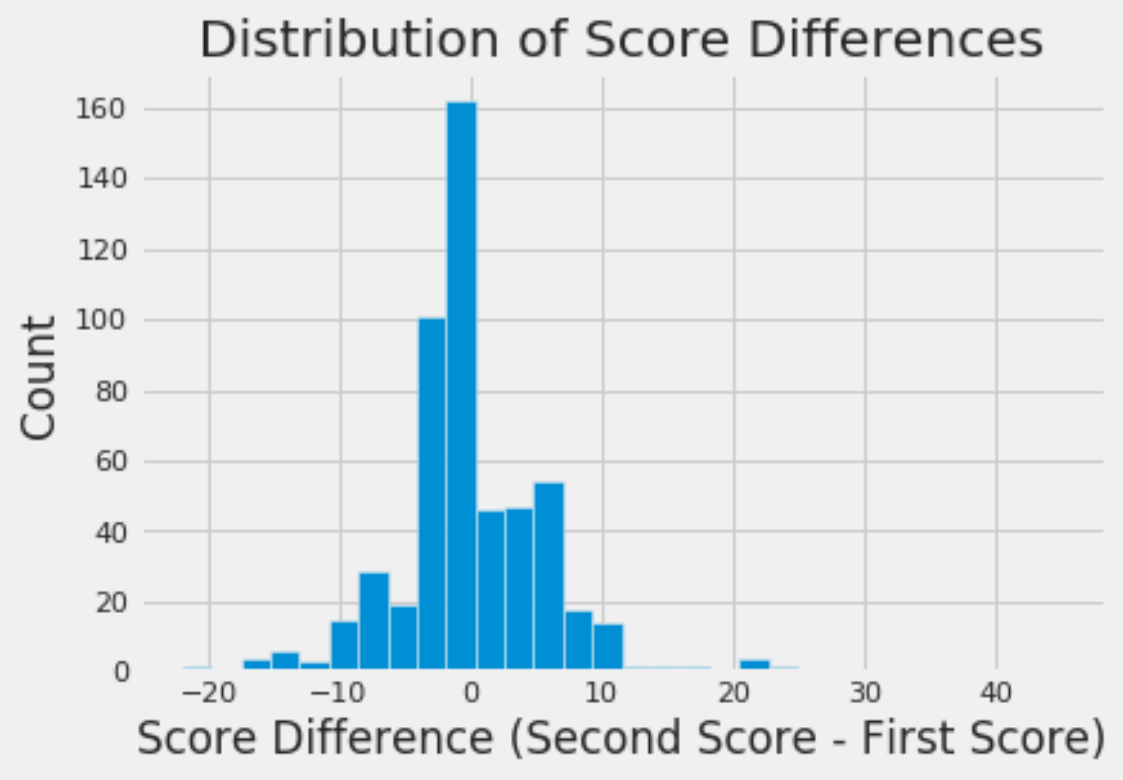
Distribution of Score Differences
Count
160
140
120
100
80
60
40
20
0
−20
−10
0
10
20
30
40
Score Difference (Second Score - First Score)

0.1.3 Question 7e

If restaurants' scores tend to improve from the first to the second inspection, what do you expect to see in the scatter plot that you made in question 7c? What do you oberve from the plot? Are your observations consistent with your expectations?

Hint: What does the slope represent?

From the scatter plot we can see that the graph has a positive relationship between the First and Second Scores. The red line is the best fit line that is computed with all the blue points. It looks like most of the data is clustered around 100 for both the first and second scores.

0.1.4 Question 7f

If a restaurant's score improves from the first to the second inspection, how would this be reflected in the histogram of the difference in the scores that you made in question 7d? What do you oberve from the plot? Are your observations consistent with your expectations? Explain your observations in the language of Statistics: for instance, the center, the spread, the deviation etc.

Yes the idea behind the histogram is that your second score should be higher than the first. This would result in a bigger remainder. If difference is negative then that means the second inspection was worst than the first. my observations are consistent with my expectation becuase most of the data was clusered in the high 80s to 100 and so I would assume that the score would be better the second time around. I think the one part im confused on is why the highest count is a negative vlaue on the difference scale however, I would assume that the original score was very close to the high 90s.

0.1.5 Question 7g

To wrap up our analysis of the restaurant ratings over time, one final metric we will be looking at is the distribution of restaurant scores over time. Create a side-by-side boxplot that shows the distribution of these scores for each different risk category from 2017 to 2019. Use a figure size of at least 12 by 8.

The boxplot should look similar to the sample below. Make sure the boxes are in the correct order!

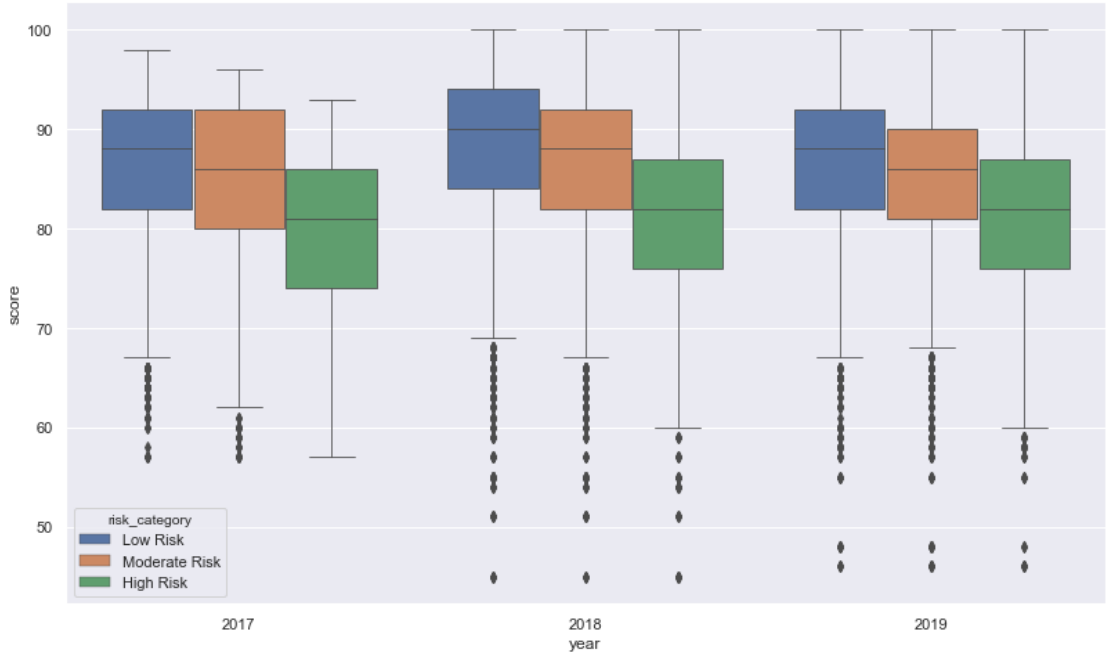

Hint: Use `sns.boxplot()`. Try taking a look at the first several parameters. The documentation is linked here!

Hint: Use `plt.figure()` to adjust the figure size of your plot.

```
In [89]: # Do not modify this line
         sns.set()
         plt.figure(figsize = (12,8))
         desmerge = vio.merge(ins2vio, how='left', on='vid')
         insmerge = ins.merge(desmerge, on='iid')
         merge = insmerge[insmerge['year']>= 2017]

         sns.boxplot(x='year', y='score', data = merge, hue='risk_category', hue_order=['Low Risk', "Mo
```

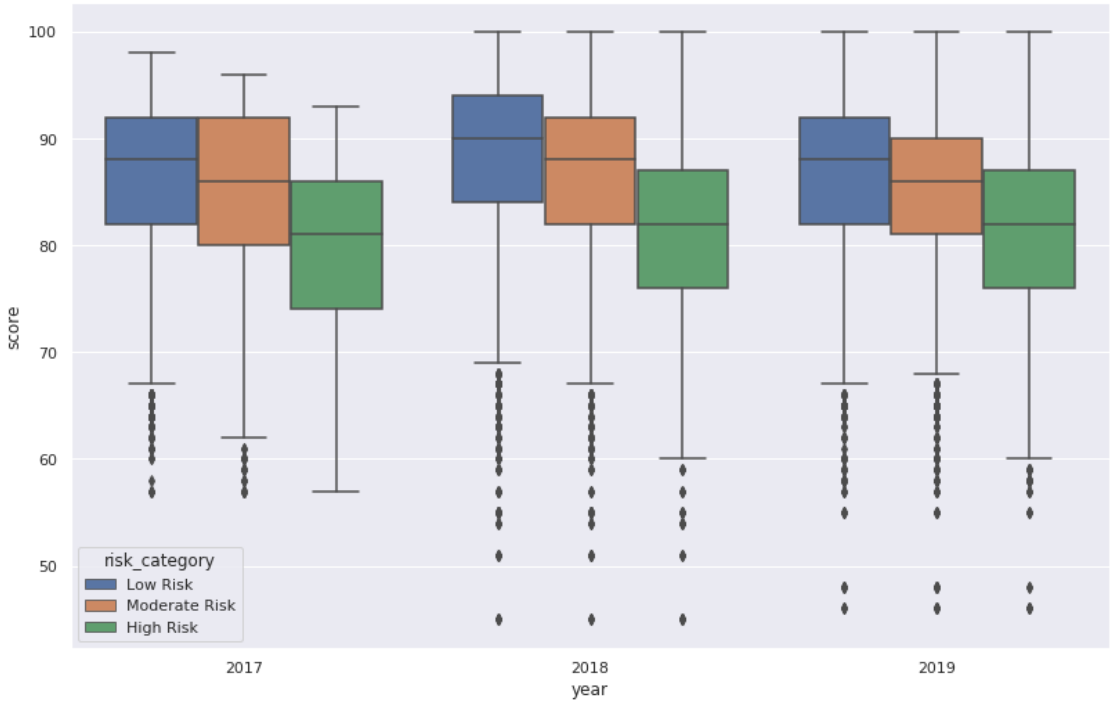
100
90
80
70
60
50
score
risk_category
Low Risk
Moderate Risk
High Risk
2017
2018
2019
year

1 8: Open Ended Question

1.1 Question 8a

1.1.1 Compute Something Interesting

Play with the data and try to compute something interesting about the data. Please try to use at least one of groupby, pivot, or merge (or all of the above).

Please show your work in the cell below and describe in words what you found in the same cell. This question will be graded leniently but good solutions may be used to create future homework problems.

1.1.2 Grading

Since the question is more open ended, we will have a more relaxed rubric, classifying your answers into the following three categories:

- **Great** (4 points): Uses a combination of pandas operations (such as groupby, pivot, merge) to answer a relevant question about the data. The text description provides a reasonable interpretation of the result.
- **Passing** (1-3 points): Computation is flawed or very simple. The text description is incomplete but makes some sense.
- **Unsatisfactory** (0 points): No computation is performed, or a computation with completely wrong results.

Please have both your code and your explanation in the same one cell below. Any work in any other cell will not be graded.

```
In [112]: q8a = vio.drop(columns=['vid']).groupby(by='risk_category').count()
          q8a['% total'] = (q8a['description']/sum(q8a['description']))*100
          q8a

          #ins.iloc[0]

          # The first observation that I see from the output below is the inspections have less than
          # half with low risk category score. I also noticed that if you add the high risk and moderat
          # risk you get majority of the data (24.61+32.30 = 56.81%). From the data analytic standpoint
          # this is concerning that most inspections were barely passing/high risk for many violations.
```

Out[112]:

risk_category	description	% total
High Risk	16	24.615385
Low Risk	28	43.076923
Moderate Risk	21	32.307692

1.1.3 Grading

Since the question is more open ended, we will have a more relaxed rubric, classifying your answers into the following three categories:

- **Great** (4 points): The chart is well designed, and the data computation is correct. The text written articulates a reasonable metric and correctly describes the relevant insight and answer to the question you are interested in.
- **Passing** (1-3 points): A chart is produced but with some flaws such as bad encoding. The text written is incomplete but makes some sense.
- **Unsatisfactory** (0 points): No chart is created, or a chart with completely wrong results.

We will lean towards being generous with the grading. We might also either discuss in discussion or post on Piazza some examplar analysis you have done (with your permission)!

You should have the following in your answers: * a few visualizations; Please limit your visualizations to 5 plots. * a few sentences (not too long please!)

Please note that you will only receive support in OH and Piazza for Matplotlib and seaborn questions. However, you may use some other Python libraries to help you create you visualizations. If you do so, make sure it is compatible with the PDF export (e.g., Plotly does not create PDFs properly, which we need for Gradescope).

In [ ]:
```
0,70,"Poor"
71,85,"Needs Improvement"
86,90,"Adequate"
91,100,"Good"
```

In [177]:
```
x_ins= ins['year']
y_ins = ins['score']
plt.hist(y_ins)
#plt.axis([2016, 2019, 0, 100])
plt.xlabel('Score')
plt.ylabel('Count')
plt.title('Count of Scores');
# This is a great representation that majority of the inspections that take place are
# in good standing. They seem to stay in the upper 80s to 100. I think what is interesting
# about this graph is that there are a good amount of restaurants
```

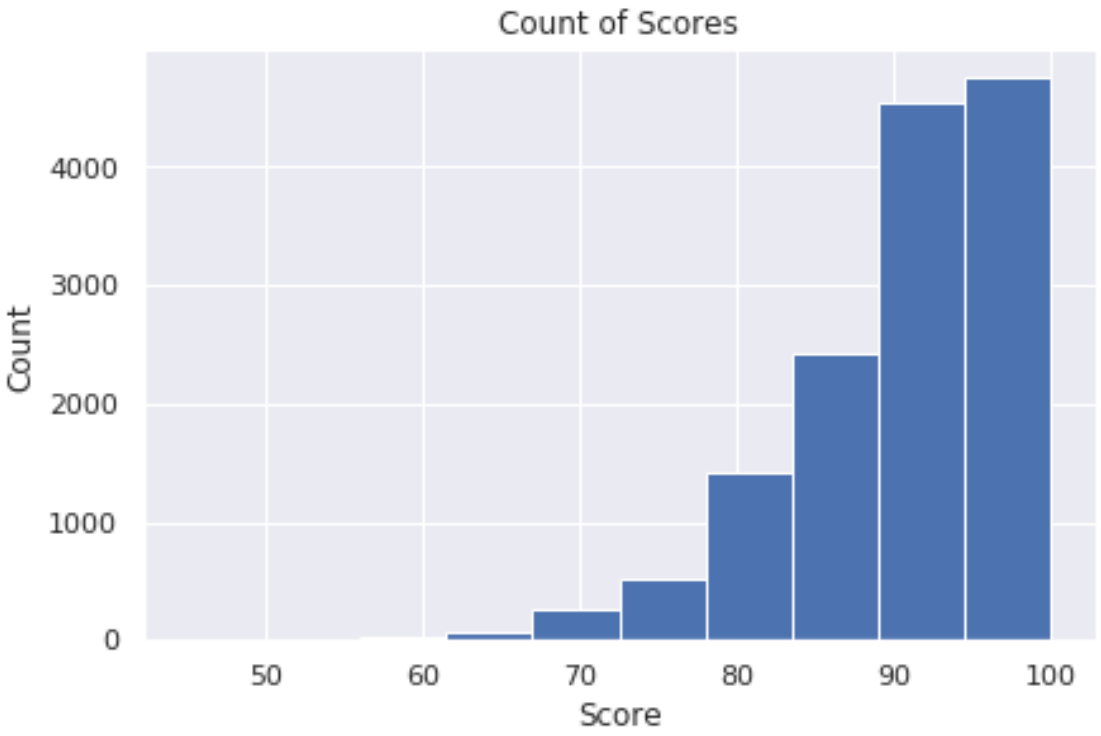
Count of Scores
Count
4000
3000
2000
1000
0
50
60
70
80
90
100
Score